

# PLATE SPINNING: MODERN CONCURRENCY in PYTHON 3

ThoughtWorks®



ThoughtWorks®

# LUCIANO RAMALHO

---

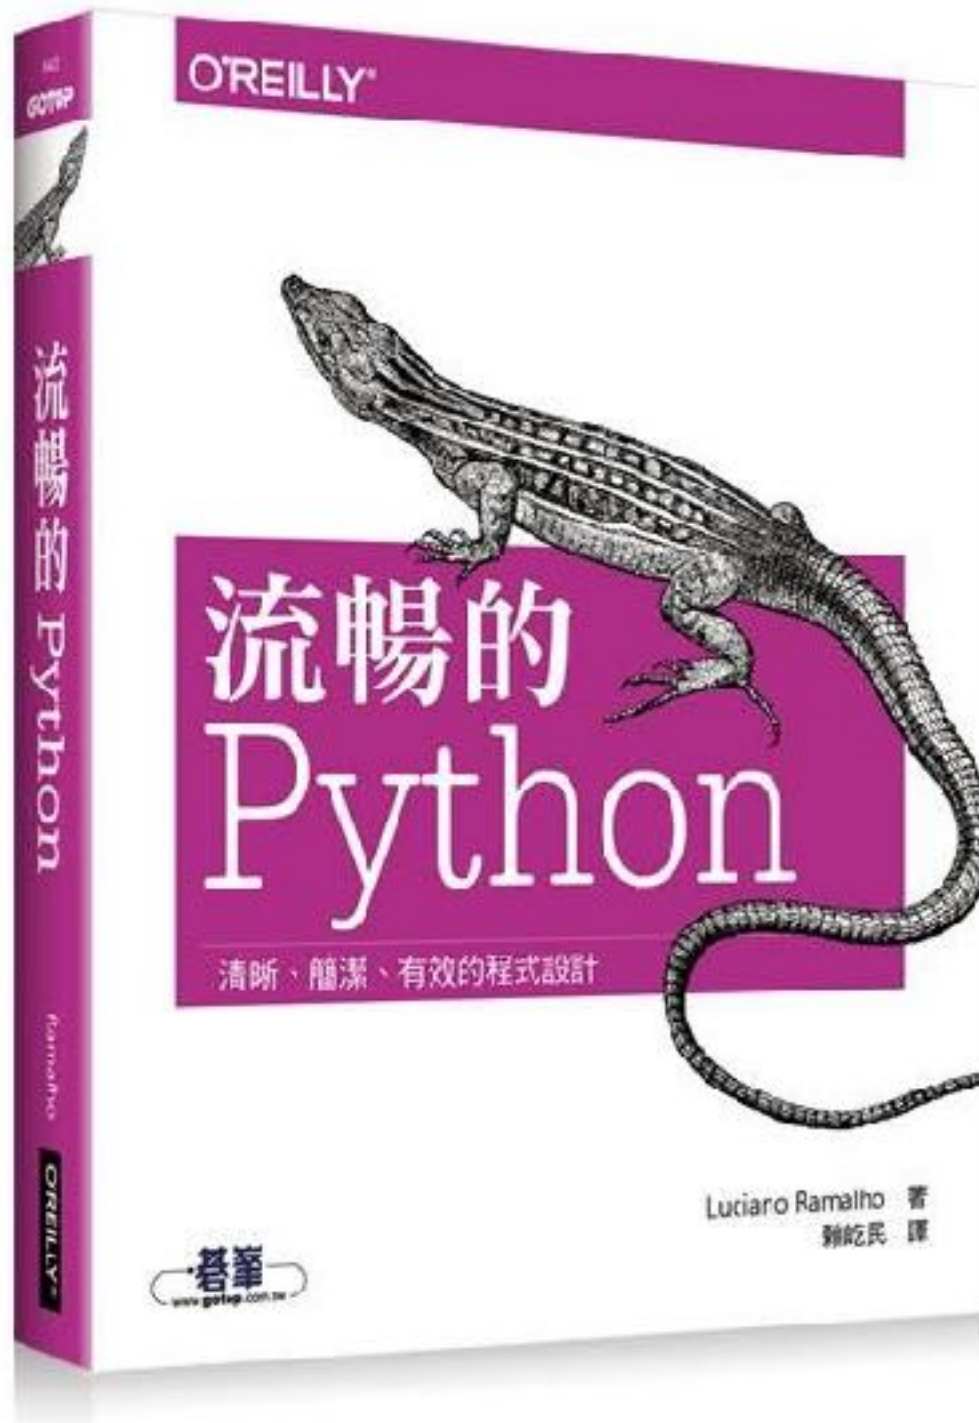
*Technical Principal*

---

@ramalhoorg  
luciano.ramalho@thoughtworks.com

# FLUENT PYTHON, MY FIRST BOOK

---



★★★★★ 4.7 stars at Amazon.com

**Fluent Python** (O'Reilly, 2015)

**Python Fluente** (Novatec, 2015)

**Python к вершинам мастерства\*** (DMK, 2015)

**流暢的 Python<sup>†</sup>** (Gotop, 2016)

also in **Polish, Korean...**

\* *Python. To the heights of excellence*

† *Smooth Python*

# CONCURRENCY

---

Not the same as parallelism



# PLATE SPINNING

---

The essential idea of concurrency: spinning 18 plates does not require 18 hands.

You can do it with 2 hands, if you know when each plate needs an intervention to keep spinning.



# CONCURRENCY VS. PARALLELISM

---

## Concurrency vs. parallelism

Concurrency is about dealing with lots of things at once.

Parallelism is about doing lots of things at once.

Not the same, but related.

Concurrency is about structure, parallelism is about execution.

Concurrency provides a way to structure a solution to solve a problem that may (but not necessarily) be parallelizable.



*Rob Pike - 'Concurrency Is Not Parallelism'*  
[https://www.youtube.com/watch?v=cN\\_DpYBzKso](https://www.youtube.com/watch?v=cN_DpYBzKso)

## EXERCISE 2

---

Creating a façade for the Unicode signs server

# CONCURRENCY DESPITE THE GIL

---

The price of the Global Interpreter Lock



# PYTHON ALTERNATIVES

---

- **Threads:**

OK for high performance I/O on constrained settings

- See: Motor 0.7 Beta With Pymongo 2.9 And A Threaded Core  
— A. Jesse Jiryu Davis — <https://emptysqua.re/blog/motor-0-7-beta/>

- **GIL-releasing threads:**

*some* external libraries in Cython, C, C++, FORTRAN...

- **Multiprocessing:** multiple instances of Python

- **Celery** and other distributed task queues

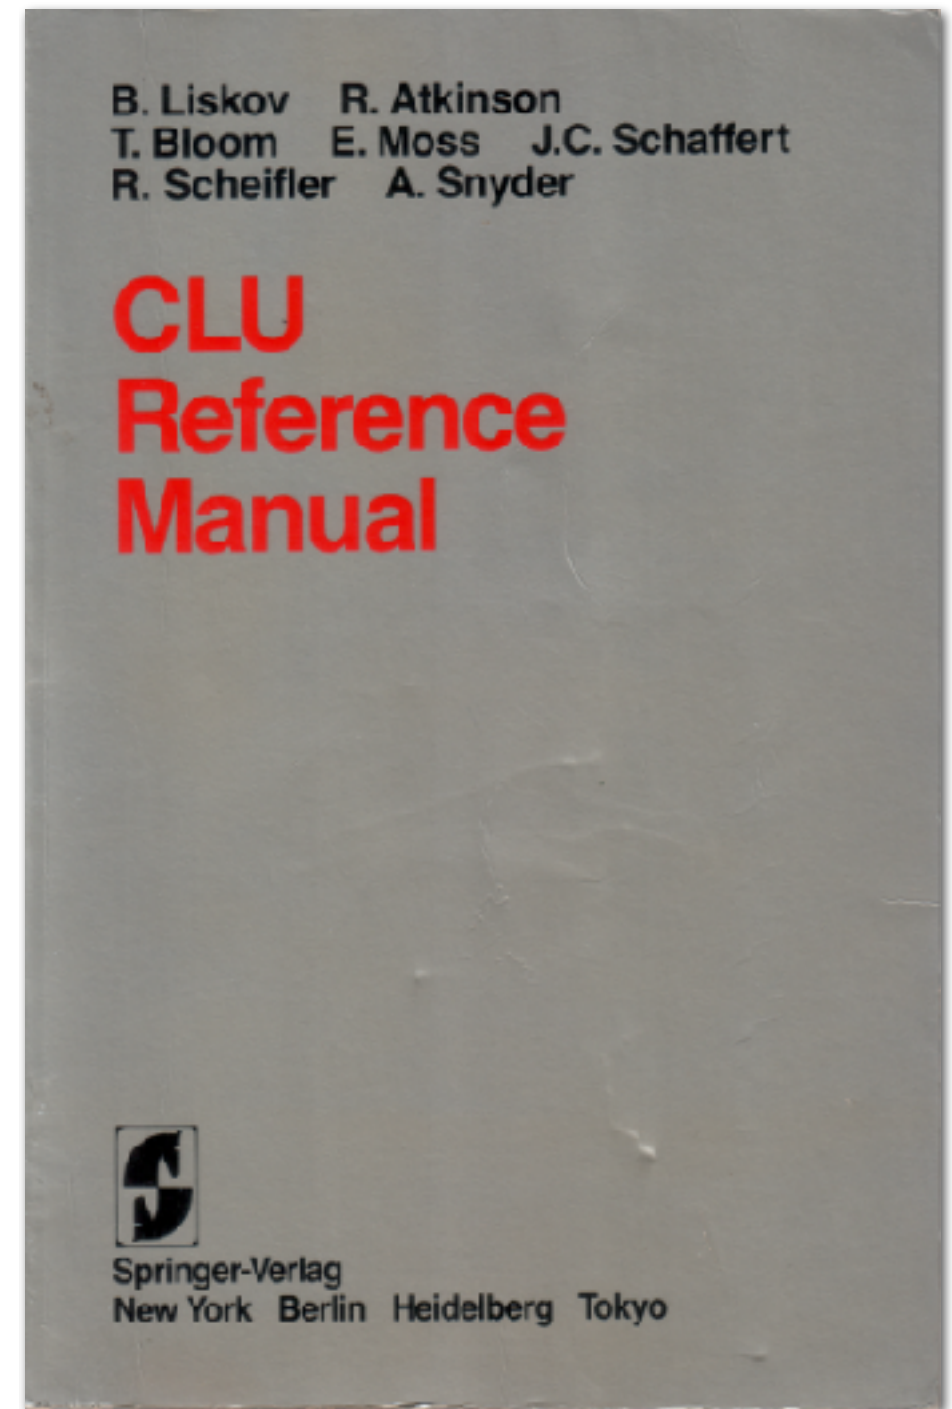
- Callbacks & deferreds in **Twisted**

- **gevent**: greenlets with monkey-patched libraries

- Generators and coroutines in **Tornado** e **Asyncio**

# GENERATORS WITH YIELD: WORK BY BARBARA LISKOV

© 2010 Kenneth C. Zirkel — CC-BY-SA

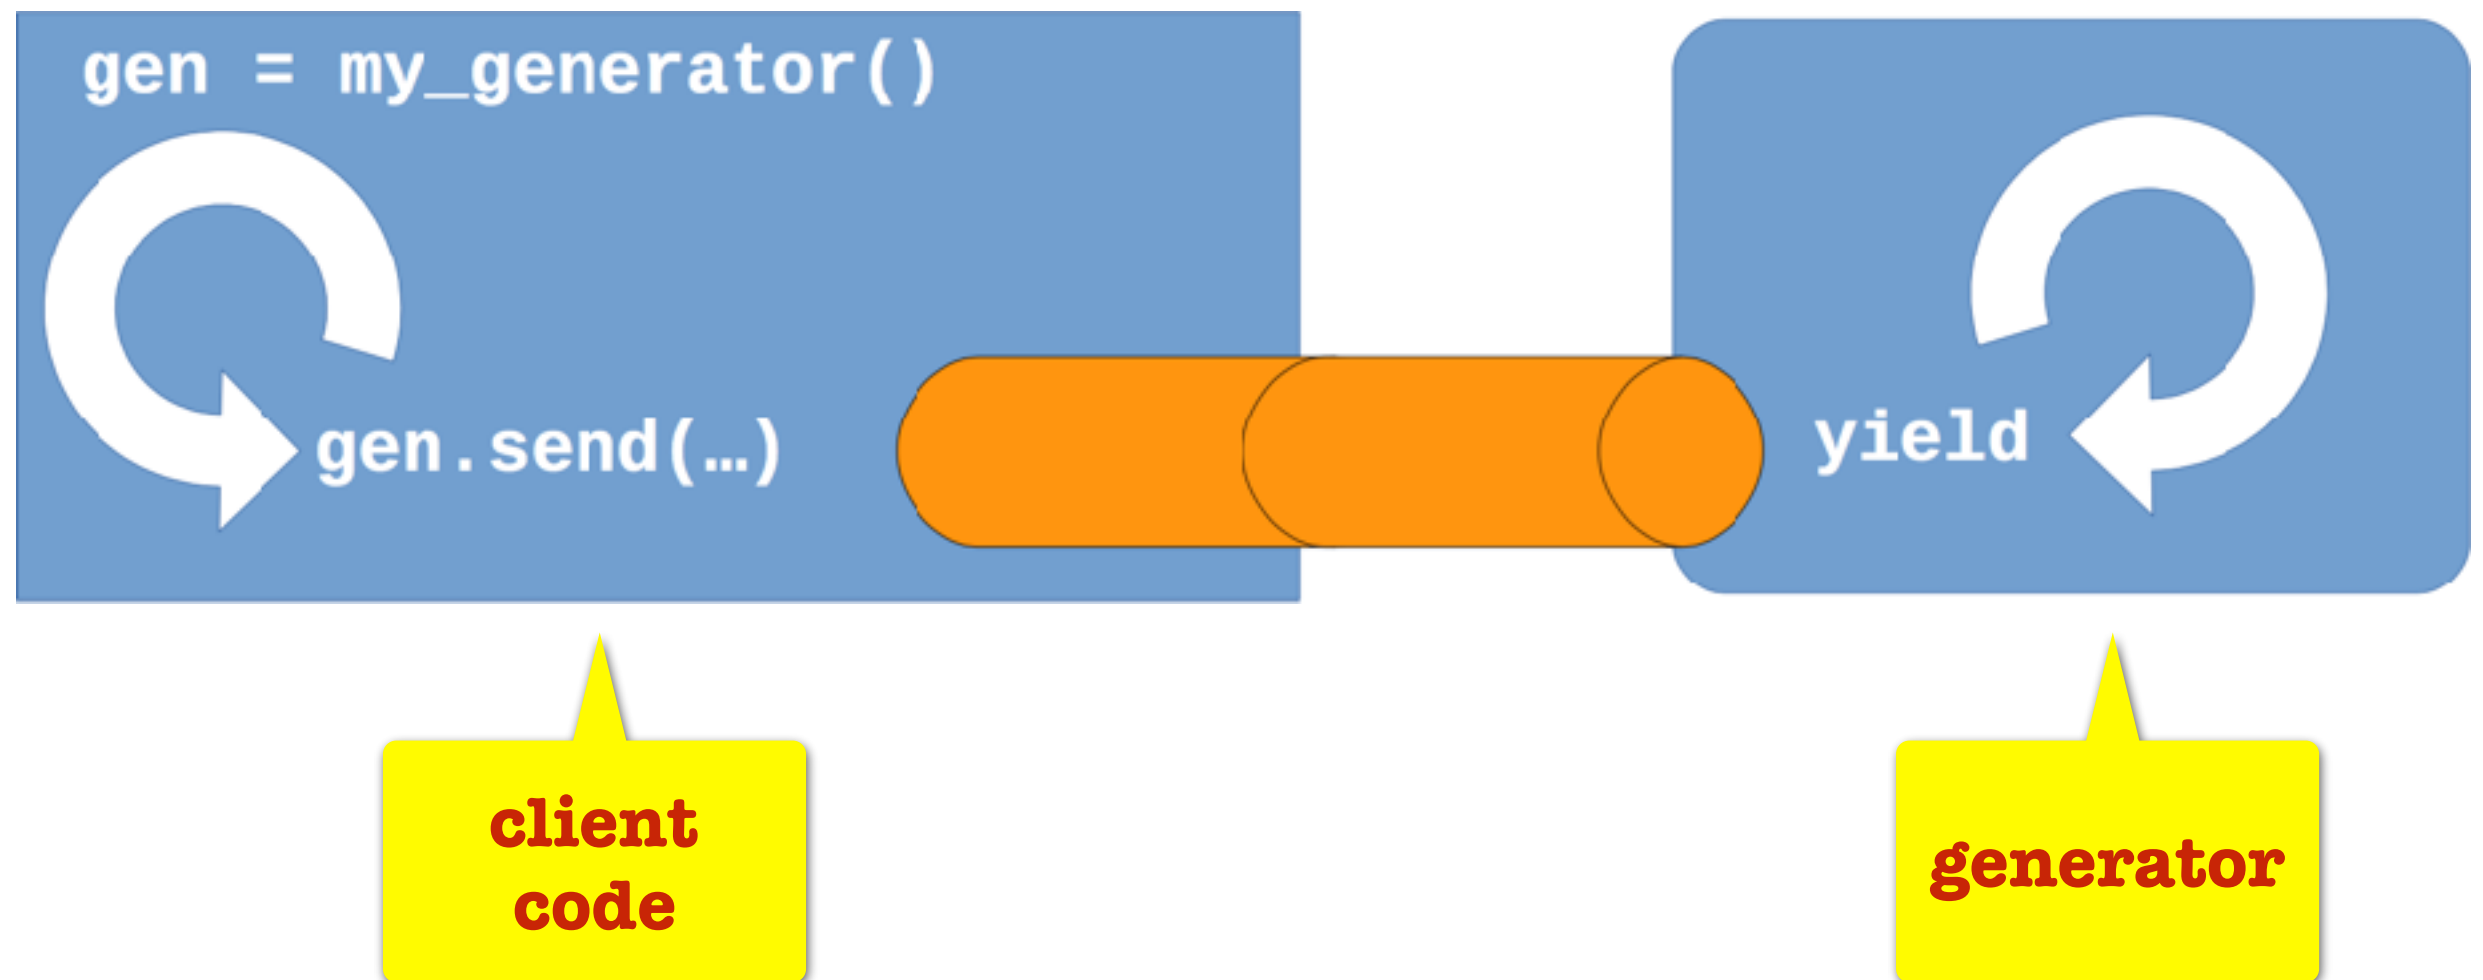


**CLU Reference Manual** — B. Liskov et. al. — © 1981 Springer-Verlag — also available online from MIT:  
<http://publications.csail.mit.edu/lcs/pubs/pdf/MIT-LCS-TR-225.pdf>

# CONCURRENCY WITH COROUTINES (1)

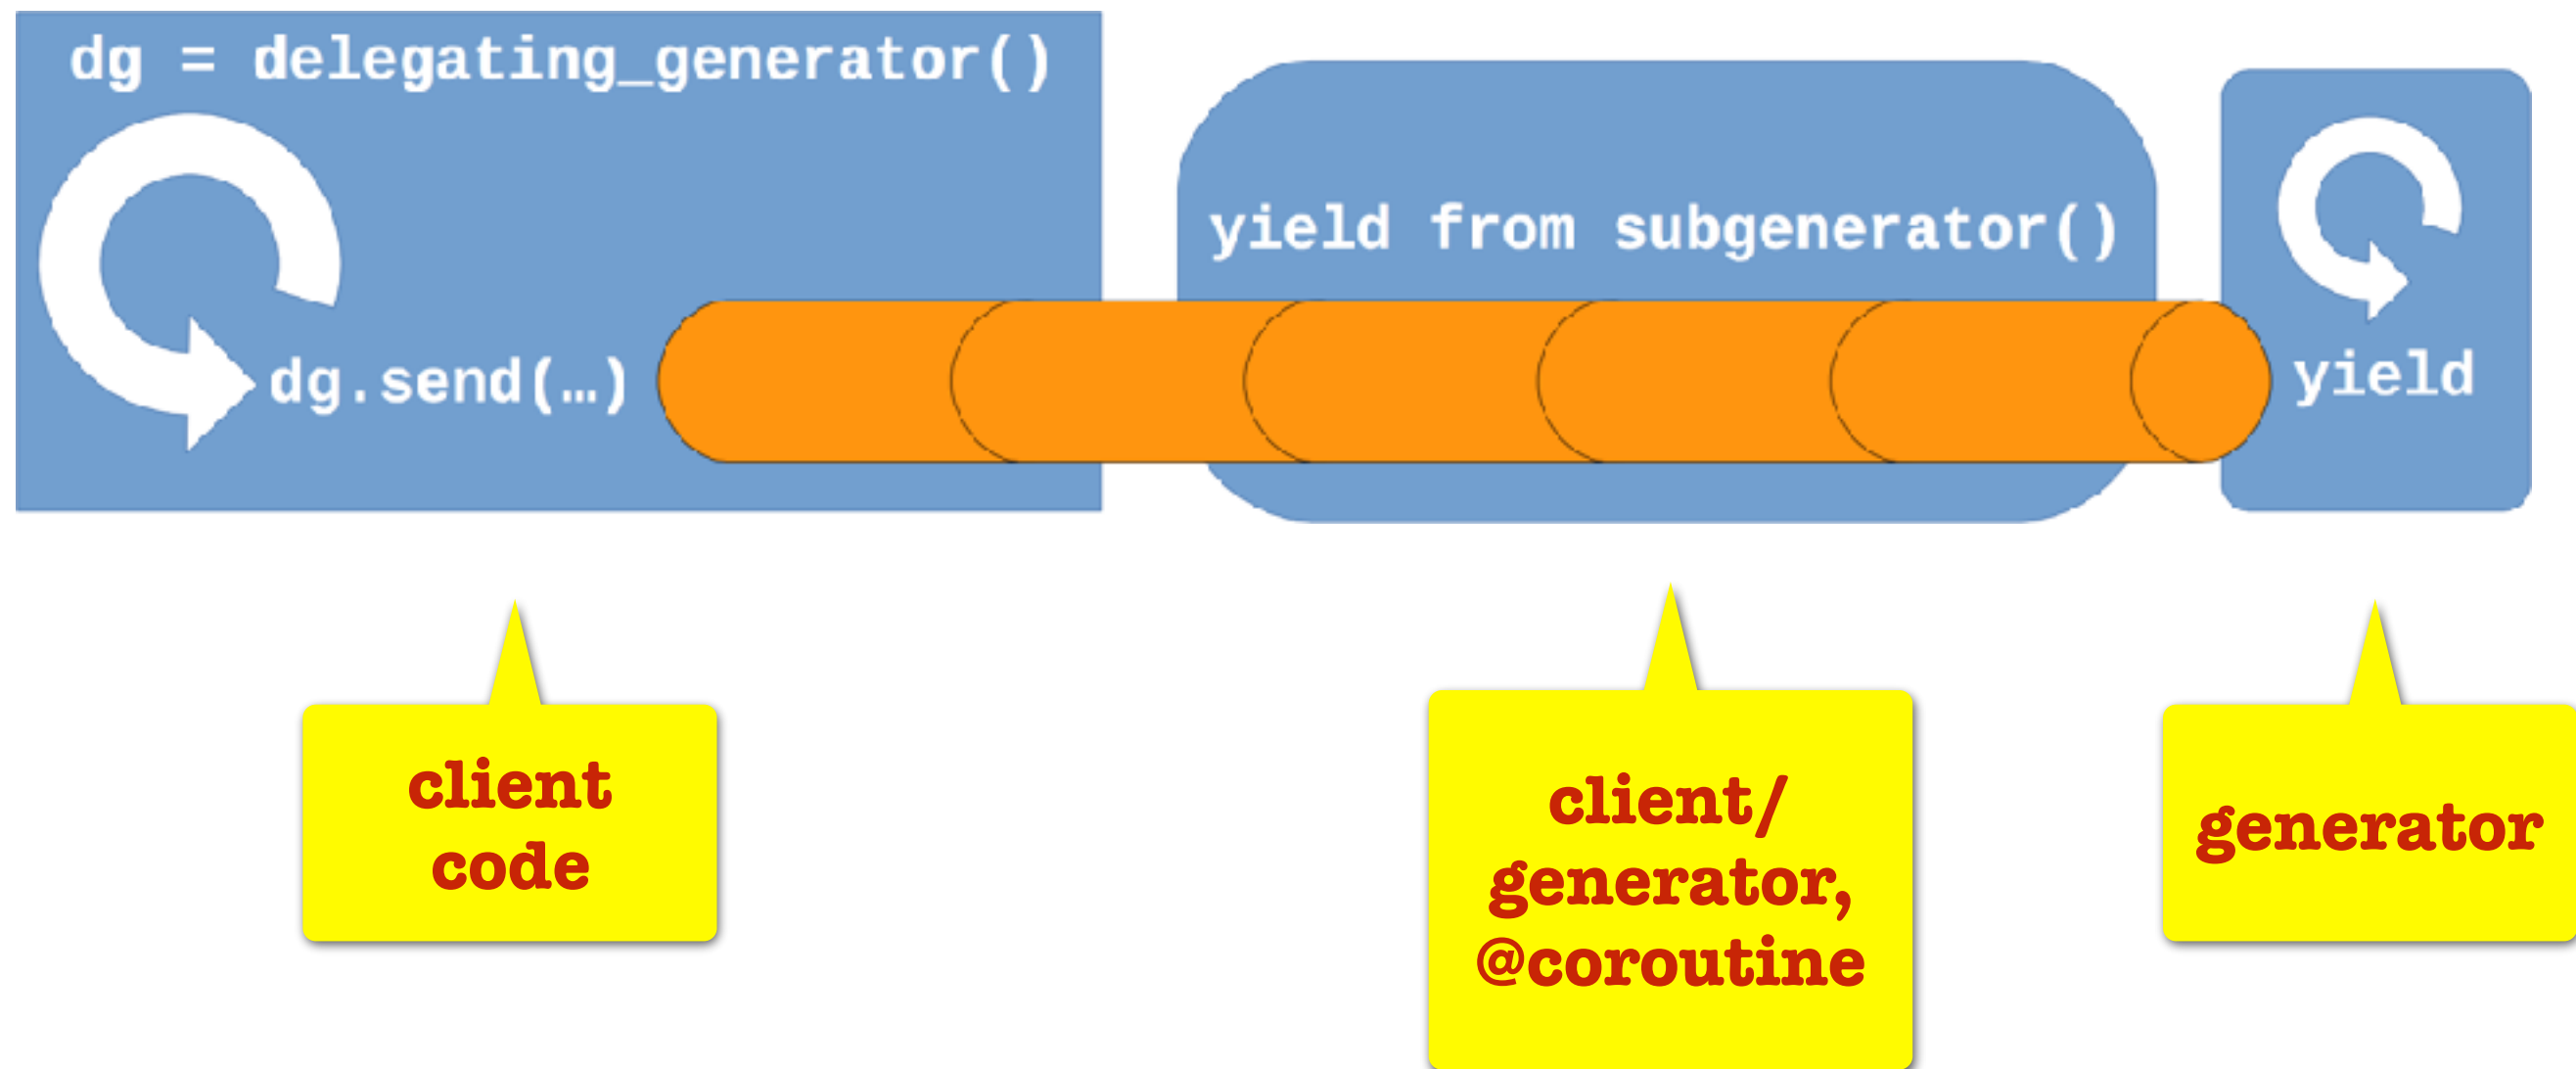
---

- In Python 2.5 (2006), the modest **generator** was enhanced with a **.send()** method



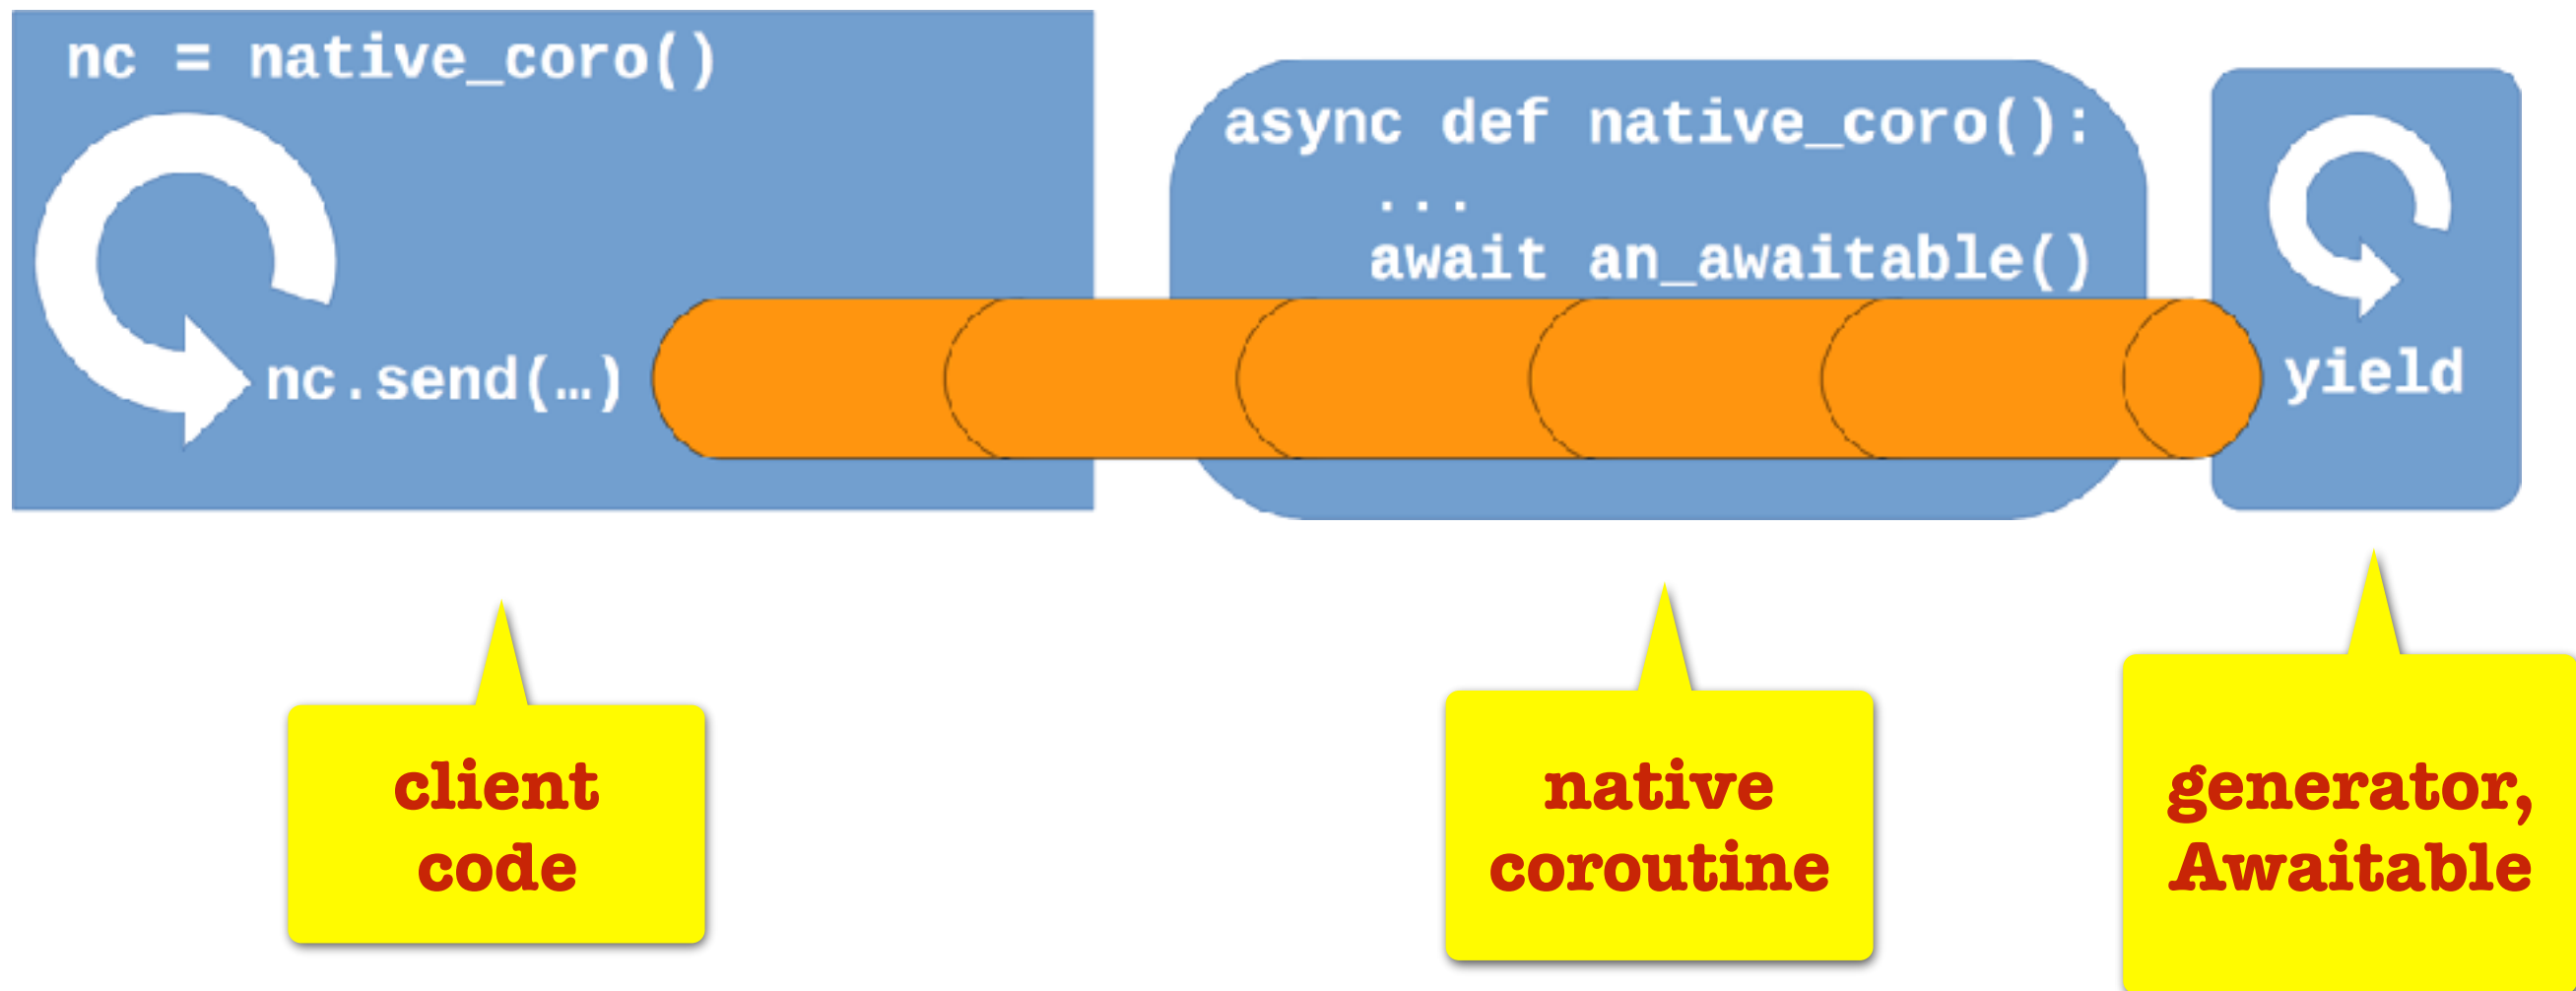
## CONCURRENCY WITH COROUTINES (2)

- In Python 3.3 (2012), the **yield from** syntax allowed a generator to delegate to another generator...



# CONCURRENCY WITH COROUTINES (3)

- Finally, in Python 3.5 (2015), **native coroutines** were born



# ASYNCH/AWAIT

---

Where the action is today



# ThoughtWorks®

# ASYNCIO

---

First use case for yield from

# ASYNC/AWAIT IS NOT JUST FOR ASYNCIO

---

- In addition to **asyncio**, there are (at least) **curio** and **trio** leveraging native coroutines for asynchronous I/O with very different APIs.
- Brett Cannon's launchpad.py example: native coroutines with a toy event loop in 120 lines using only the packages **time**, **datetime**, **heapq** and **types**.

# EXERCISE 1

---

The countdown experiment

# ASYNCIO: FIRST PACKAGE TO LEVERAGE ASYNC/AWAIT

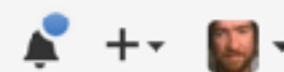
---

- Package designed by Guido van Rossum (originally: Tulip)
  - added to Python 3.4, provisional status up to Python 3.5: significant API changes
- **asyncio** is no longer provisional in *Python 3.6*
  - most of the API is rather low-level: support for library writers
  - no support for HTTP in the standard library: aiohttp is the most cited add-on
- Very active eco-system
  - see: <https://github.com/aio-libs/>



This organization Search

Pull requests Issues Gist



## aio-libs

The set of asyncio-based libraries built with high quality for humans

<https://groups.google.com/forum/#!forum/aio-libs>

Repositories

People 8

Filters

Find a repository...

### aiomysql

Python ★ 197 29

aiomysql is a library for accessing a MySQL database from the asyncio

Updated a day ago

### aiohttp\_admin

Python ★ 24 4

admin interface for aiohttp application

Updated 2 days ago

### aiokafka

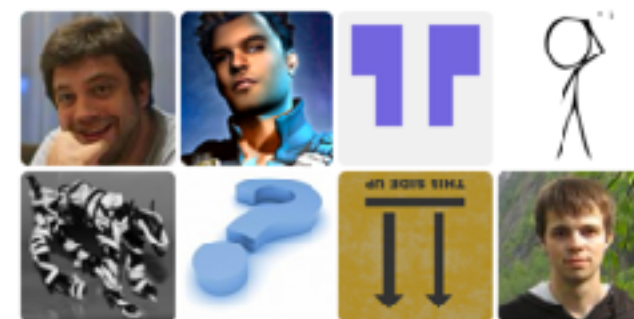
Python ★ 49 12

asyncio client for kafka

Updated 2 days ago

### People

8 >





## aiozmq

Python ★ 175 🔗 23

Asyncio (pep 3156) integration with ZeroMQ

Updated 5 days ago

## aiopg

Python ★ 332 🔗 48

aiopg is a library for accessing a PostgreSQL database from the asyncio

Updated 5 days ago

## sockjs

Python ★ 53 🔗 10

SockJS Server

Updated 5 days ago

## multidict

Python ★ 13 🔗 3

multidict implementation

Updated 5 days ago

## aiobotocore

Python ★ 52 🔗 14

asyncio support for boto3 library using aiohttp

Updated 5 days ago

## aiohttp-debugtoolbar

JavaScript ★ 54 🔗 16



**aiodbc**

Python ★ 39 🔒 0

aicodbc - is a library for accessing a ODEC databases from the asyncio

Updated 2 days ago

**aioredis**

Python ★ 196 🔒 41

asyncio (PEP 3156) Redis support

Updated 2 days ago

**aiohttp\_session**

Python ★ 39 🔒 24

Provide sessions for aiohttp.web

Updated 3 days ago

**aiohttp\_jinja2**

Python ★ 49 🔒 19

jinja2 template renderer for aiohttp.web

Updated 3 days ago

**yarl**

Python ★ 33 🔒 4

Yet another URL library

Updated 4 days ago

**aiozmq**

Python ★ 175 🔒 22



## [aiohttp\\_debugtoolbar](#)

JavaScript ★ 54 🔗 16

aiohttp\_debugtoolbar is library for debugtoolbar support for aiohttp

Updated 7 days ago

## [aiohttp\\_cors](#)

Python ★ 21 🔗 6

CORS support for aiohttp

Updated 7 days ago

## [janus](#)

Python ★ 31 🔗 1

Thread-safe asyncio-aware queue

Updated 7 days ago

## [aiomcache](#)

Python ★ 35 🔗 7

Minimal asyncio memcached client

Updated 7 days ago

## [aiorwlock](#)

Python ★ 17 🔗 1

Synchronization primitive RLock for asyncio (PEP 3156)

Updated 7 days ago

Minimal asyncio memcached client  
Updated 7 days ago

**aiorwlock**

Python ★ 17 🔗 1

Synchronization primitive RLock for asyncio (PEP 3156)  
Updated 7 days ago

**aiosmtpd**

Python ★ 20 🔗 5

A reimplementation of the Python stdlib smtpd.py based on asyncio.  
Updated 8 days ago

**aiohttp\_mako**

Python ★ 9 🔗 2

mako template renderer for aiohttp.web  
Updated 8 days ago



## aiocouchdb

Python ★ 32 🔒 8

CouchDB client built on top of aiohttp (asyncio)

Updated 25 days ago

## async\_timeout

Python ★ 5 🔒 1

asyncio-compatible timeout class

Updated 28 days ago

## sphinxcontrib-asyncio

Python ★ 4 🔒 1

Sphinx extension to add asyncio-specific markups

Updated on Apr 15

## aiopyspp

Python ★ 9 🔒 2

IETF PPSP RFC7574 in Python/asyncio

Updated on Feb 3

## aiohttp

Python ★ 76 🔒 9

REST interface for server based on aiohttp (abandoned)

Updated on Apr 30, 2015

# PLUGGABLE EVENT LOOP

---

- **asyncio** includes an event loop
- The **AbstractEventLoopPolicy** API lets us replace the default loop with another implementing **AbstractEventLoop**
  - **AsyncIOMainLoop** implemented by the **Tornado** project
  - An event loop for GUI programming: **Quamash** (PyQt4, PyQt5, PySide)
  - Event loops wrapping the **libuv** library, the highly efficient asynchronous core of Node.js

# UVLOOP

---

- Implemented as Cython bindings for **libuv**
- Written by Yuri Selivanov, who proposed the **async/await** syntax
  - PEP 492 — Coroutines with async and await syntax



# USING UVLOOP

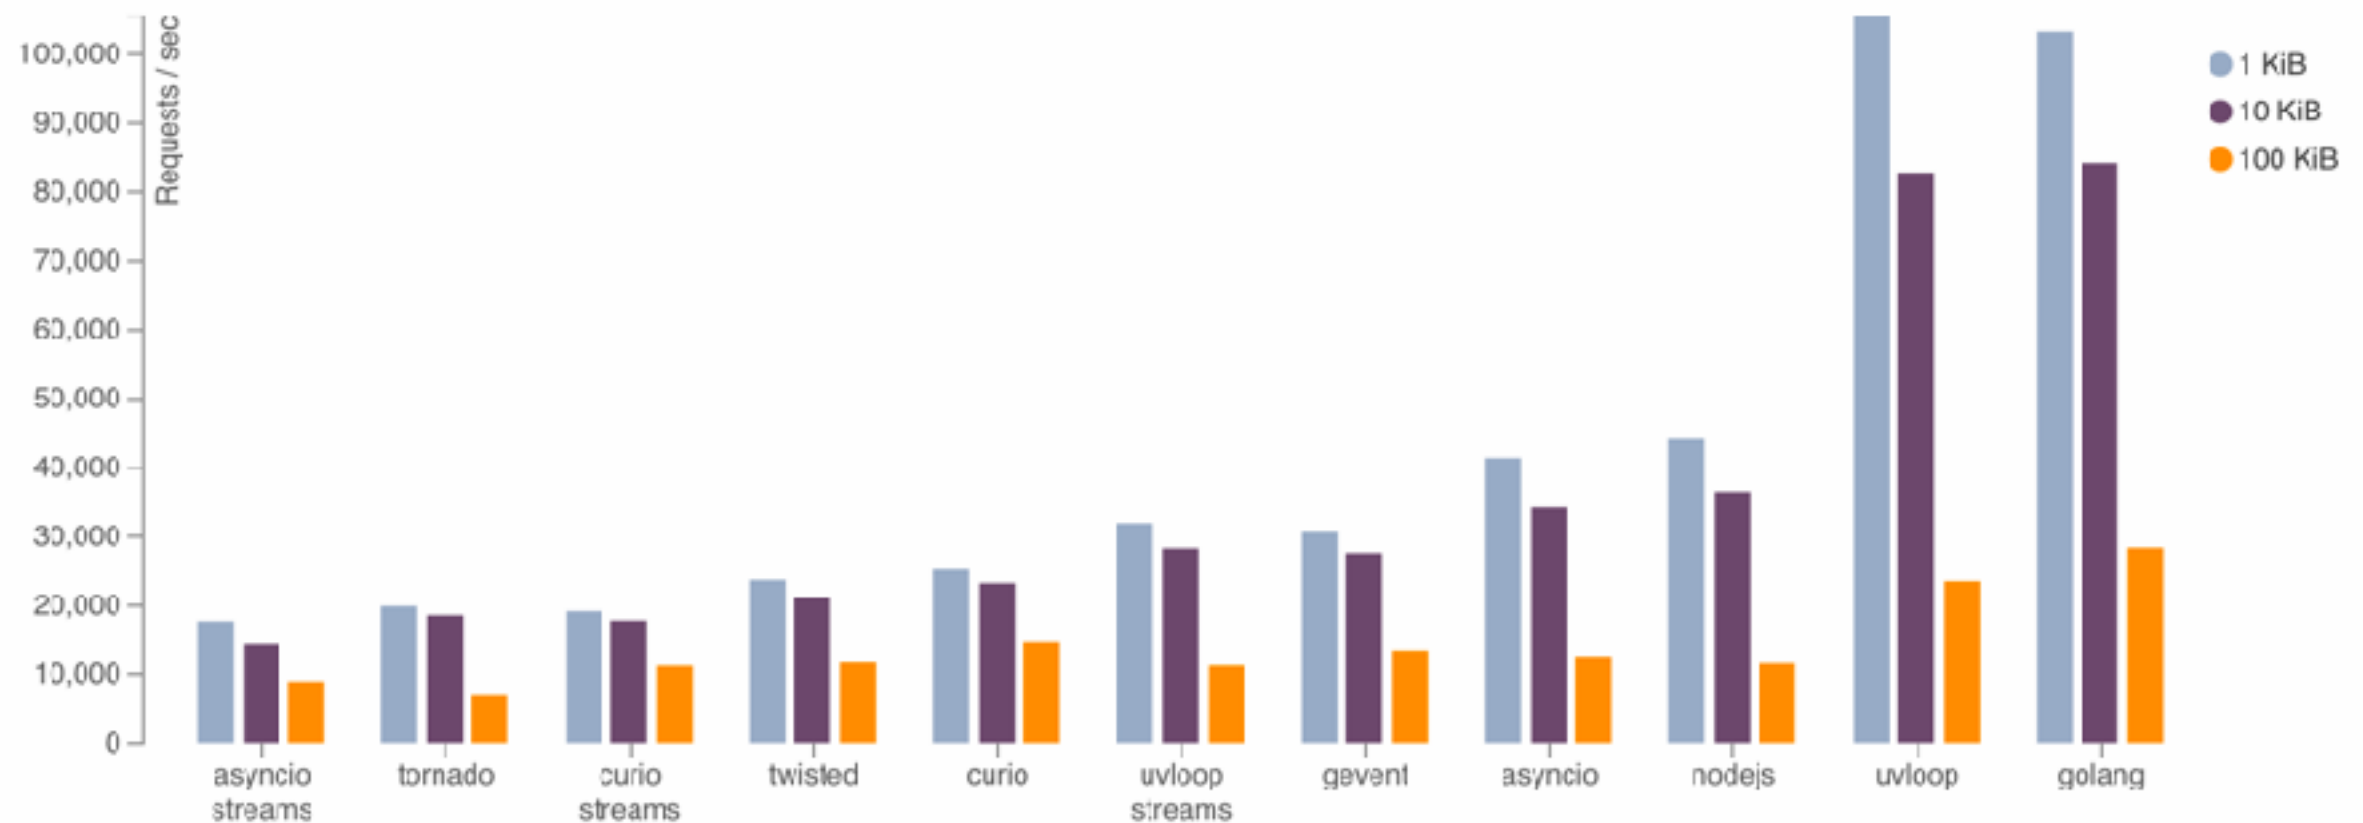
uvloop example · ramalho/t... +	
GitHub, Inc. (US)   https://github.com/ramalho/tudo-agora/commit/e1abc22e4e8de8178e2d0e45c3c70c13071301e6   vs code source repo	
... @@ -1,6 +1,6 @@	
1	1
2	2
3	3
4	4
5	5
6	6
@@ -15,6 +15,7 @@	
15	15
16	16
17	17
	18
18	19
19	20
20	21
@@ -49,7 +50,8 @@ def save_flag(img, filename):	
49	50
50	51
51	52
52	53
	54
53	55
54	56
55	57

# UVLOOP PERFORMANCE

## TCP

This benchmark tests the performance of a simple echo server with different message sizes. We use 1, 10, and 100 KiB packages. The concurrency level is 10. Each benchmark was run for 30 seconds.

See also the full TCP benchmarks [report](#).



Fonte: **uvloop: Blazing fast Python networking** — Yury Selivanov — 2016-05-03

<https://magic.io/blog/uvloop-make-python-networking-great-again/>

# NATIVE COROUTINES

---

Better syntax for asynchronous programming

# THE NEW ASYNC DEF SYNTAX

---

- **PEP 492**: New keywords introduced in Python 3.5
- **async def** to define *native coroutines*
- **await** to delegate processing to **Awaitable** objects
  - can only be used in native coroutines
- **Awaitable** or "*Future-like*":
  - Instances of **asyncio.Future** (or **Task**, a subclass of **Future**)
  - native coroutines (**async def...**)
  - generator-coroutines decorated with **@types.coroutine**
  - objects implementing **\_\_await\_\_** (which returns an iterator)

# SEQUENTIAL VS. ASYNCHRONOUS (1)

flags\_await with coroutines ...

GitHub, Inc. (US) | https://github.com/ramalho/tudo-agora/commit/32d92b02883a0b2ade14a715ce137d1e739581b2 | Pesquisar

27 countries/flags\_await.py View

... @@ -1,10 +1,10 @@

1 """Download flags of top 20 countries by population

2

3 -Sequential version

4

5 Sample run::

6

7 - \$ python3 flags\_seq.py

8 BD BR CD CN DE EG ET FR ID IN IR JP MX NG PH PK RU TR US VN

9 20 flags downloaded in 10.16s

10

@@ -13,7 +13,8 @@

13 import time

14 import sys

15

16 -import requests

17

18 POP20\_CC = ('CN IN US ID BR PK NG BD RU JP '

19 'MX PH VN ET EG DE IR TR CD FR').split()

20

@@ -34,23 +35,27 @@ def save\_flag(img, filename):

34 fp.write(img)

35

36

37

38 -def get\_flag(cc):

39 url = '{}/{cc}/{cc}.gif'.format(BASE\_URL, cc=cc.lower())

40 - resp = requests.get(url)

41

1 """Download flags of top 20 countries by population

2

3 +Asynchronous version

4

5 Sample run::

6

7 + \$ python3 flags\_await.py

8 BD BR CD CN DE EG ET FR ID IN IR JP MX NG PH PK RU TR US VN

9 20 flags downloaded in 10.16s

10

@@ -13,7 +13,8 @@

13 import time

14 import sys

15

16 +import asyncio

17 +import aiohttp

18

19 POP20\_CC = ('CN IN US ID BR PK NG BD RU JP '

20 'MX PH VN ET EG DE IR TR CD FR').split()

21

@@ -34,23 +35,27 @@ def save\_flag(img, filename):

35 fp.write(img)

36

37

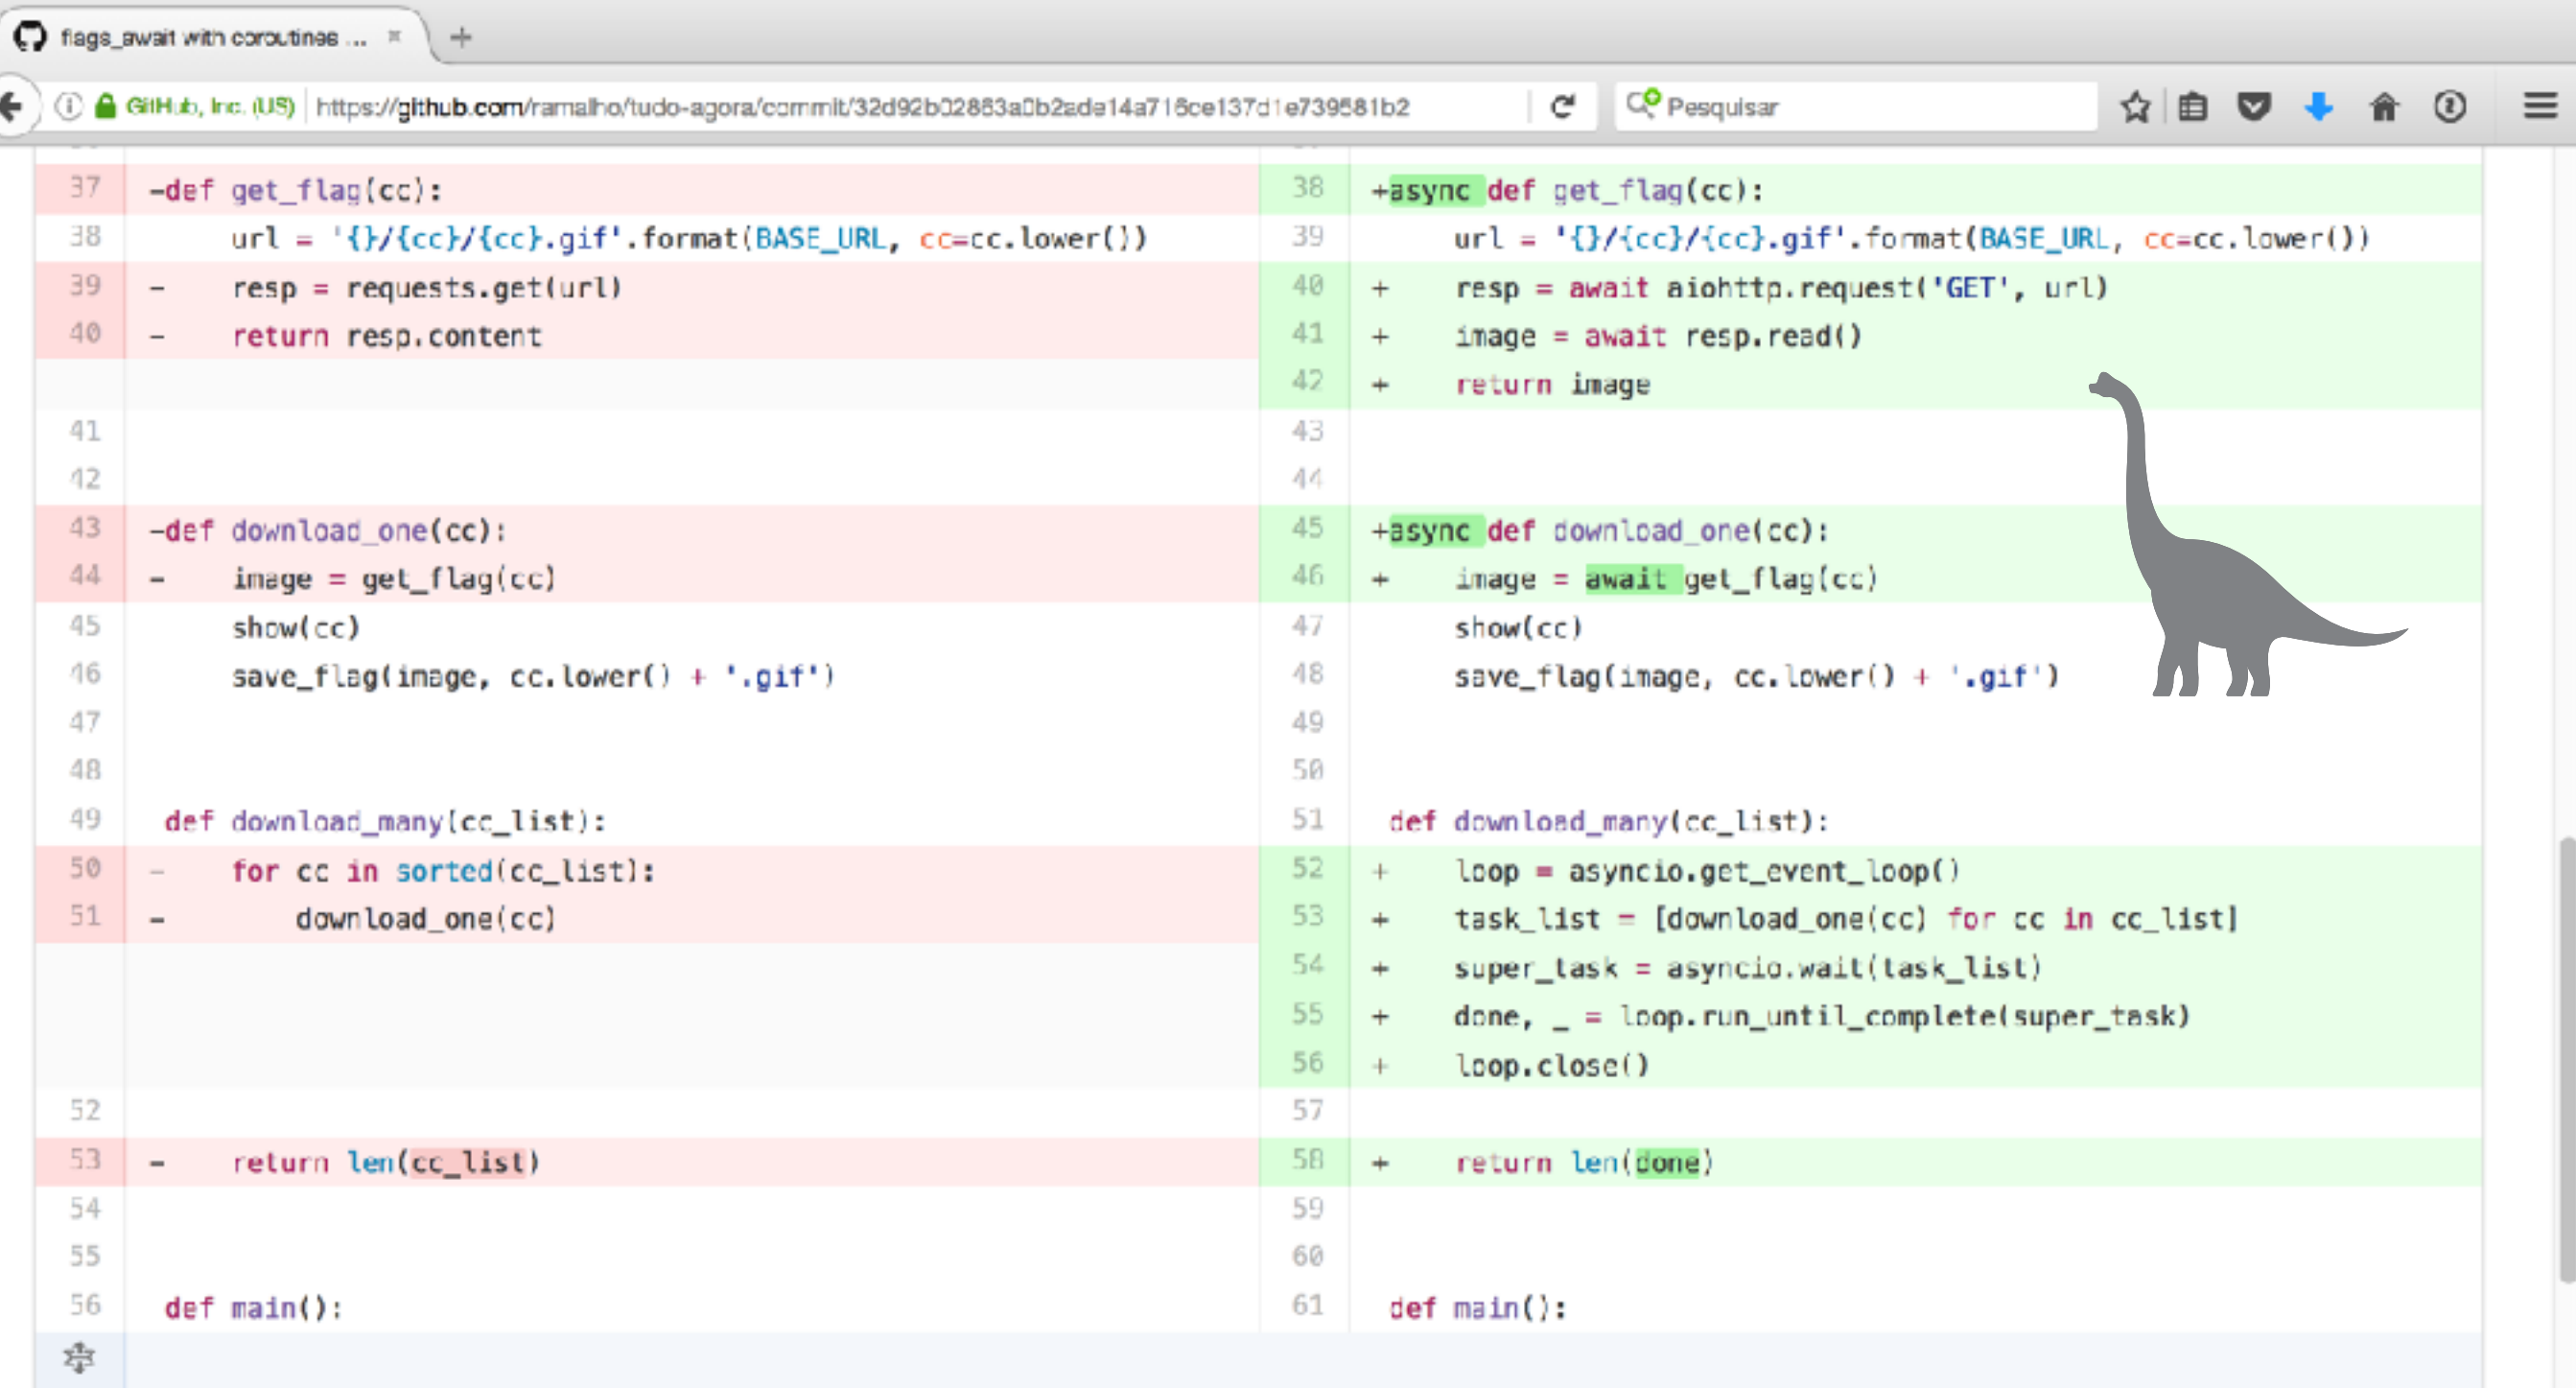
38 +async def get\_flag(cc):

39 url = '{}/{cc}/{cc}.gif'.format(BASE\_URL, cc=cc.lower())

40 + resp = await aiohttp.request('GET', url)

41

# SEQUENTIAL VS. ASYNCHRONOUS (2)



The screenshot shows a GitHub commit diff for a file named `flags_await_with_coroutines.py`. The diff compares two versions of the code: the original (left, lines 37-56) and the updated version (right, lines 38-61). The updated version uses asynchronous programming with `asyncio` and `await` to make the code more efficient.

```
37 -def get_flag(cc):  
38     url = '{}/{cc}/{cc}.gif'.format(BASE_URL, cc=cc.lower())  
39 -     resp = requests.get(url)  
40 -     return resp.content  
  
41  
42  
43 -def download_one(cc):  
44 -     image = get_flag(cc)  
45     show(cc)  
46     save_flag(image, cc.lower() + '.gif')  
47  
48  
49     def download_many(cc_list):  
50 -         for cc in sorted(cc_list):  
51 -             download_one(cc)  
  
52  
53 -     return len(cc_list)  
54  
55  
56     def main():  
57  
58  
59  
60  
61
```

The updated code (right) is as follows:

```
38 +async def get_flag(cc):  
39     url = '{}/{cc}/{cc}.gif'.format(BASE_URL, cc=cc.lower())  
40 +     resp = await aiohttp.request('GET', url)  
41 +     image = await resp.read()  
42 +     return image  
  
43  
44  
45 +async def download_one(cc):  
46 +     image = await get_flag(cc)  
47     show(cc)  
48     save_flag(image, cc.lower() + '.gif')  
49  
50  
51     def download_many(cc_list):  
52 +         loop = asyncio.get_event_loop()  
53 +         task_list = [download_one(cc) for cc in cc_list]  
54 +         super_task = asyncio.wait(task_list)  
55 +         done, _ = loop.run_until_complete(super_task)  
56 +         loop.close()  
  
57  
58 +     return len(done)  
59  
60  
61     def main():
```

A cartoon dinosaur is drawn on the right side of the diff, standing next to the updated code.

0 comments on commit 32d92b0

Lock conversation



Write

Preview

AA B i

“ < > ↺

≡ ≡ ≡

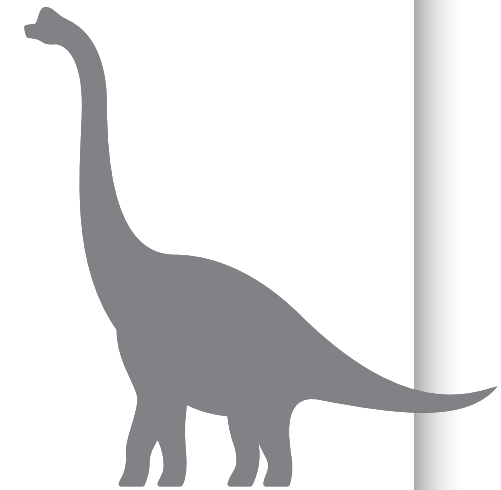
↶ @ ★



# NATIVE COROUTINES IN ACTION

---

```
38  async def get_flag(cc):
39      url = '{}/{cc}/{cc}.gif'.format(BASE_URL, cc=cc.lower())
40      resp = await aiohttp.request('GET', url)
41      image = await resp.read()
42      return image
43
44
45  async def download_one(cc):
46      image = await get_flag(cc)
47      show(cc)
48      save_flag(image, cc.lower() + '.gif')
49
50
51  def download_many(cc_list):
52      loop = asyncio.get_event_loop()
53      task_list = [download_one(cc) for cc in cc_list]
54      super_task = asyncio.wait(task_list)
55      done, _ = loop.run_until_complete(super_task)
56      loop.close()
57
58      return len(done)
```



# MORE SYNTACTIC SUPPORT

---

- **PEP 492** also introduced:
  - **async with:**  
invokes asynchronous special methods **`__aenter__*`** and **`__aexit__*`**
    - **`*`**: coroutines (return **`Awaitable`** objects)
  - **async for:**  
invokes special methods **`__aiter__`** e **`__anext__*`**
    - **`__aiter__`**: not a coroutine, but returns an asynchronous iterator
    - asynchronous integrator implements **`__anext__*`** as a coroutine

# EXAMPLE USING ASYNC WITH AND ASYNC FOR

```
1  import asyncio
2  import aiopg
3
4  dsn = 'dbname=aiopg user=aiopg password=passwd host=127.0.0.1'
5
6  async def go():
7      async with aiopg.create_pool(dsn) as pool:
8          async with pool.acquire() as conn:
9              async with conn.cursor() as cur:
10                 await cur.execute("SELECT 1")
11                 ret = []
12                 async for row in cur:
13                     ret.append(row)
14                 assert ret == [(1,)]
15
16  loop = asyncio.get_event_loop()
17  loop.run_until_complete(go())
```

# STILL MORE SYNTACTIC SUPPORT

---

- New features in **Python 3.6**:
  - **PEP 525**: Asynchronous Generators (!)
  - **PEP 530**: Asynchronous Comprehensions

# WRAPPING UP

---

The end is near

# MY TAKE ON ASYNCIO

---

- Young ecosystem: libraries evolving fast
  - even trivial examples in Fluent Python now issue warnings or are broken
- **asyncio** with is open for better implementation thanks to its pluggable event loop policy
  - alternative event loops available for a while:
    - Tornado: **AsyncIOMainLoop**
    - QT: **Quamash**
    - libuv: **uvloop** and **pyuv**
- Give **Python 3.6** a try before jumping to Go, Elixir or Node

# THE ONE (ABSTRACT) LOOP

---

*One Loop to rule them all, One Loop to find them,  
One Loop to bring them all and in liveness bind them.*

# FACEBOOK: PYTHON IN PRODUCTION ENGINEERING

FirefoxArquivoEditarExibirHistóricoFavoritosFerramentasJanelaAjuda

Python in production engine... +

https://code.facebook.com/posts/1040181199381023/python-in-production-engineering/

vs code source repo

Search

AndroidiOSWebBackendHardware

27 de maioPRODUCTION ENGINEERING · PYTHON

## Python in production engineering

Romain Komorn

Python aficionados are often surprised to learn that Python has long been the language most commonly used by **production engineers** at Facebook and is the third most popular language at Facebook, behind Hack (our in-house dialect of PHP) and C++. Our engineers build and maintain thousands of Python libraries and binaries deployed across our entire infrastructure.

Every day, dozens of Facebook engineers commit code to Python utilities and services with a wide variety of purposes including binary distribution, hardware imaging, operational automation, and infrastructure management.

### Python at Facebook by the numbers

### Recommended

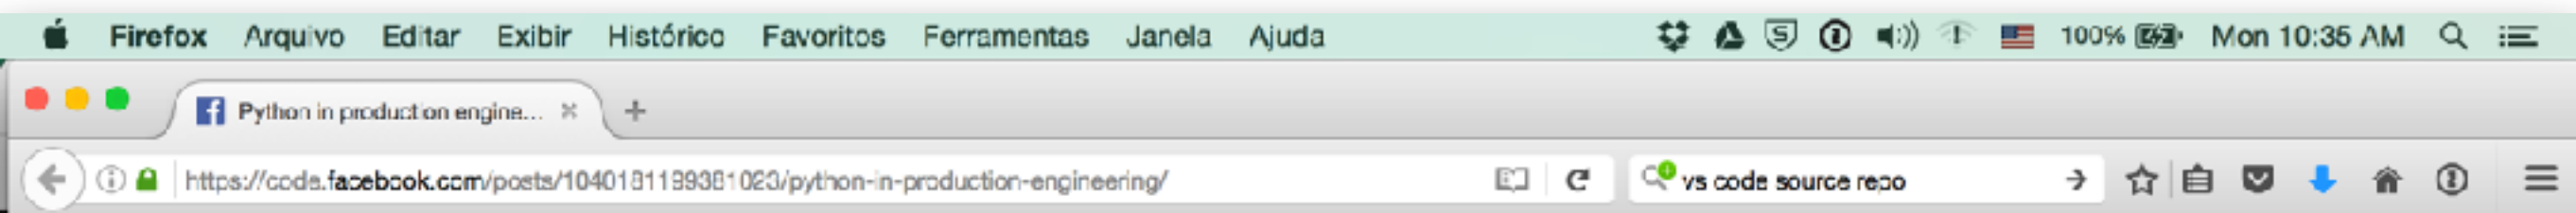
Type A

Type B

Serving Facebook Multifeed: Efficiency, performance gains through redesign

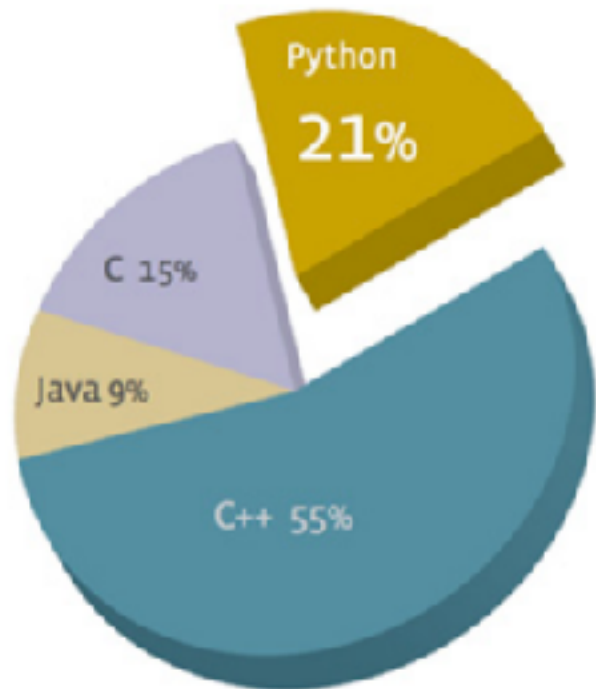


# PYTHON AT FACEBOOK



## Python at Facebook by the numbers

- 21 percent of Facebook Infrastructure's codebase



- Millions of lines of code, thousands of libraries and binaries
- 2016 to date: average 5,000 commits per month, 1,000+ committers
- 5 percent Py3 (as of May 2016)

## Python in production engineering

Python is heavily used by the Facebook infrastructure teams and is ubiquitous in production engineering. Teams typically maintain Python client libraries (generally Thrift) for their services,

**SECURITY @ SCALE**  
Security @Scale 2015:  
Engineering Security

```
-----BEGIN RSA PRIVATE KEY-----
MIIEpAIBAAKCAQEA...
> ssh
Scalable and secure access
with SSH
```

Join Optimization in Apache  
Hive

# PYTHON 3 + ASYNCIO AT FACEBOOK

Firefox Arquivo Editar Exibir Histórico Favoritos Ferramentas Janela Ajuda

Python in production engine... \*

https://code.facebook.com/posts/1040161199361023/python-in-production-engineering/

vs code source repo


## Python 3 deployments

Facebook's scale pushes Python's performance to its limits. Our codebase features various models and libraries (Twisted, Gevent, futures, AsyncIO, and many others). All ports and new projects use Python 3 unless Python 2 support is absolutely necessary. Currently, 5 percent of our Python services in production are running Python 3.

The following Python 3-compatible projects have already been open-sourced:

- [FBOSS CLI](#) — a Python 3.5 CLI that hits thrift APIs on Facebook in house switch agent
- [Facebook Python Ads API](#) — compatible with Python 3
- [FBTFTP](#) — a dynamic TFTP server framework written in Python 3
- [PYAIB](#) — Python Async IrcBot framework

There is a lot of exciting work to be done in expanding our Python 3 codebase. We are increasingly relying on AsyncIO, which was introduced in Python 3.4, and seeing huge performance gains as we move codebases away from Python 2. We hope to contribute more performance-enhancing fixes and features back to the Python community.

 Curtir

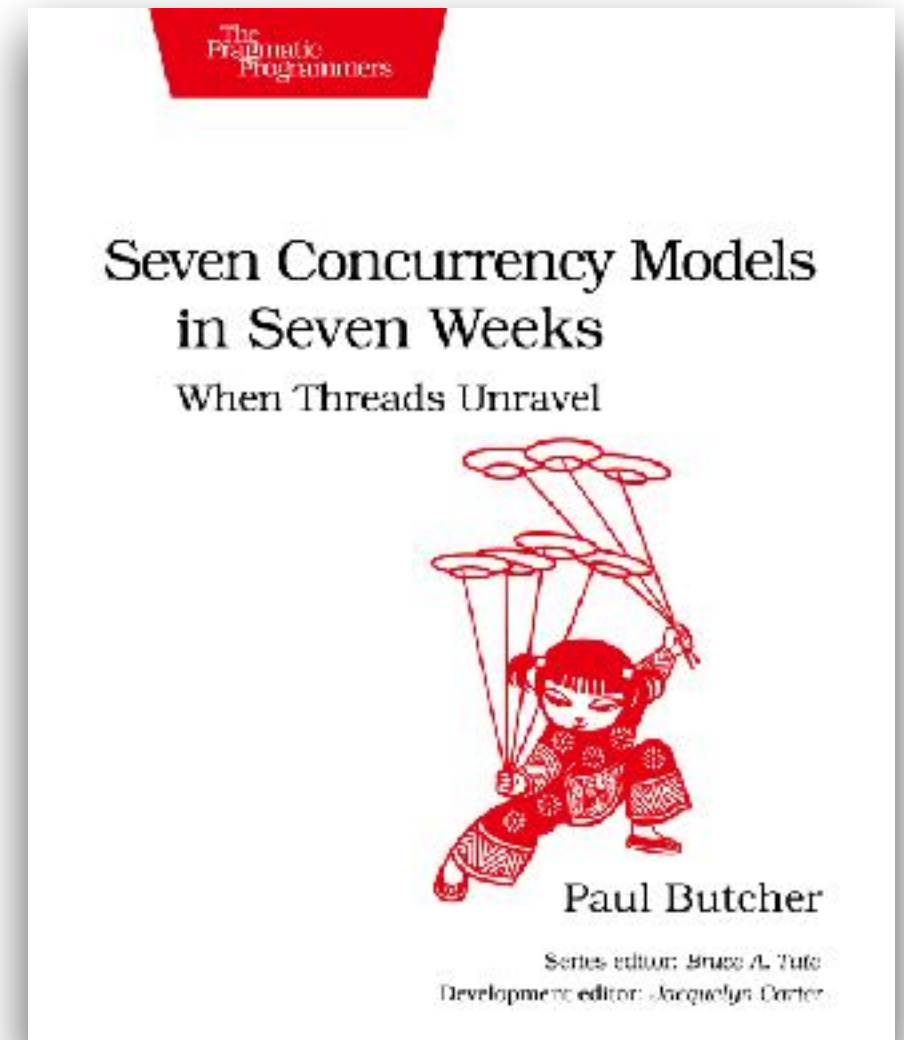
Compartilhar

1 mil pessoas curtiram isso.

# UNRAVELLING THREADS

---

- *Seven Concurrency Models in Seven Weeks — When Threads Unravel* (Paul Butcher)
- Callbacks are not covered
- Chap. 1: the problem with threads
- Remaining 6 chapters: more powerful abstractions
  - Actors, CSP, STM, data parallelism...
- Native support in languages
  - Erlang, Elixir, Clojure, Go, Cilk, Haskell...



## OTHER USES FOR ASYNC/AWAIT

---

Python's loose coupled introduction of new syntax with semantics based on `__special_methods__` allows even more experimentation than new asyncio event loops.

Libraries using async/await with very different APIs:

- David Beazley's **curio**
- Nathaniel... 's **trio**

# ASYNCIO VS. CURIO

---

The spinner example

# ¿QUESTIONS?

ThoughtWorks®

**LUCIANO RAMALHO**

---

*Technical Principal*

---

@ramalhoorg  
luciano.ramalho@thoughtworks.com

ThoughtWorks®